

Daily Stress Recognition from Mobile Phone Data, Weather Conditions and Individual Traits

Andrey Bogomolov
University of Trento,
SKIL Telecom Italia Lab
Via Sommarive, 5
I-38123 Povo - Trento, Italy
andrey.bogomolov@unitn.it

Bruno Lepri
Fondazione Bruno Kessler
Via Sommarive, 18
I-38123 Povo - Trento, Italy
lepri@fbk.eu

Michela Ferron
Fondazione Bruno Kessler
via Sommarive, 18
I-38123 Povo - Trento, Italy
ferron@fbk.eu

Fabio Pianesi
Fondazione Bruno Kessler
Via Sommarive, 18
I-38123 Povo - Trento, Italy
pianesi@fbk.eu

Alex (Sandy) Pentland
MIT Media Lab
20 Ames Street
Cambridge, MA, USA
pentland@mit.edu

ABSTRACT

Research has proven that stress reduces quality of life and causes many diseases. For this reason, several researchers devised stress detection systems based on physiological parameters. However, these systems require that obtrusive sensors are continuously carried by the user. In our paper, we propose an alternative approach providing evidence that daily stress can be reliably recognized based on behavioral metrics, derived from the user's mobile phone activity and from additional indicators, such as the weather conditions (data pertaining to transitory properties of the environment) and the personality traits (data concerning permanent dispositions of individuals). Our multifactorial statistical model, which is person-independent, obtains the accuracy score of 72.28% for a 2-class daily stress recognition problem. The model is efficient to implement for most of multimedia applications due to highly reduced low-dimensional feature space (32d). Moreover, we identify and discuss the indicators which have strong predictive power.

Categories and Subject Descriptors

I.5 [Computing Methodologies]: PATTERN RECOGNITION—*Models, Design Methodology, Implementation*; J.4 [Computer Applications]: SOCIAL AND BEHAVIORAL SCIENCES—*Sociology, Psychology*

General Terms

Algorithms; Experimentation; Measurement; Theory

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
MM'14, November 3–7, 2014, Orlando, Florida, USA.
Copyright 2014 ACM 978-1-4503-3063-3/14/11 ...\$15.00.
<http://dx.doi.org/10.1145/2647868.2654933>.

Keywords

stress recognition; mobile sensing; pervasive computing

1. INTRODUCTION

Nowadays, the number of mobile phones in use worldwide is about 5 billion, with millions of new subscribers everyday¹. Mobile phones allow for unobtrusive and cost-efficient access to huge streams of previously inaccessible data related to daily social behavior [29]. These devices are able to sense a wealth of behavioral data such as (i) location, (ii) other devices in physical proximity through Bluetooth scanning, (iii) communication data, including both metadata (logs of who, when, and duration) of phone calls and text messages (sms), etc. Correspondingly, the availability is continuously growing of huge streams of *personal data* related to activities, routines and social interactions [11, 29] which represent a novel opportunity to address fundamental problems of our societies in different fields, such as mobility and urban planning [19], finance [46], healthy living and subjective well-being [31, 33].

In this work, we focus on one of the most widespread and debilitating problem of our subjective well-being and our society: stress. Stress is a well-known condition in modern life and research has shown that the amount of cumulative stress plays a role in a broad range of physical, psychological and behavioural conditions, such as anxiety, low self-esteem, depression, social isolation, cognitive impairments, sleep and immunological disorders, neurodegenerative diseases and other medical conditions [8], while also significantly contributing to healthcare costs. Hence, measuring stress in daily life situations has become an important challenge [39]. Today, the availability of huge and diverse streams of pervasive data produced by and about people allows for automatic, unobtrusive, and fast recognition of daily stress levels. An early prediction of stress symptoms can indeed help to prevent situations that are risky for human life [23].

Several studies have produced interesting results that support the feasibility of detecting stress levels through phys-

¹<http://www.ericsson.com/ericsson-mobility-report>

iological sensors (see [23], [26]). However, the use of physiological sensors is limited by several shortcomings. Stress detection systems based on physiological measurement such as heart-rate variability or skin conductance are intrusive and need to be easily wearable to be exploited in natural settings; the data they produce can be confounded by daily life activities such as speaking or drinking; they exhibit important between-person differences [39].

Recently, social psychologist Miller wrote “The Smartphone Psychology Manifesto” in which he argued that the smartphones should be seriously considered as new research tools for psychology. In his opinion, these tools could revolutionize all fields of psychology and other behavioral sciences making these disciplines more powerful, sophisticated, and grounded in real-world behavior [36] and [30]. Indeed, several works have started to use smartphone activity data in order to detect and predict personality traits [6, 9, 37, 48], mood states [31], and daily happiness [38]. Stopczynski et al. [49] described the Copenhagen Networks Study, a large-scale study designed to measure human interactions spanning multiple years.

Smartphones data can be used to detect stress levels as well. Indeed, stress levels are associated with the type of activities people engage in, including those executed at/through their smartphone (for instance, a high number of phone calls and/or e-mails from many different people could be associated with higher stress levels). Weather conditions – an environmental transitory property – in turn, have been argued [24], [41] to be often associated with stress, acting either directly (as stressors) or indirectly (by affecting individual sensitivity to stressors). Finally, the impact of all these transitory factors – (smartphone) activities and weather conditions – on stress induction can be expected to be modulated by personal characteristics and differences [50], [52]. For example, a neurotic person could react with higher levels of stress to a high number of interactions (call, sms or proximity interactions) than an emotionally stable person; an extrovert or agreeable person, in turn, might well find him/herself at ease with a high number of interactions.

In this paper, we approach the automatic recognition of daily stress as a 2-class classification problem (non-stressed vs stressed) based on information concerning different types of data: a) people activities, as detected through their smartphones; b) weather conditions; c) personality traits. The information about people activities is represented by features extracted from call and sms logs and from Bluetooth hits, able to capture (i) the amount of calls, of sms and of proximity interactions; (ii) the diversity of calls, of sms, and of proximity interactions; and (iii) regularity in user behaviors. In addition, we use weather conditions (environmental and transitory factors) along with personality traits (internal and stable factors); the latter are mediating factors that can modulate people responses to stressors (e.g., weather, daily activity). This multifactorial approach will be compared to approaches based only on a family of features (personality, weather conditions, mobile phone features) or simpler combinations of families of features (personality and weather conditions; personality and mobile phone features; weather conditions and mobile phone features).

Classification experiments are performed using a variety of approaches and the best solution for our classification problem was found using an ensemble of tree classifiers based on a Random Forest algorithm. Our multifactorial approach

obtains an accuracy score of 72.28% for a 2-class daily stress recognition problem, providing evidence that individual daily stress can be reliably predicted from the combination of smartphone usage data, weather conditions and individual dispositions (personality traits). Interestingly, if one of these information sources is dropped, the recognition performances decrease drastically.

In sum, the main contributions of this paper are as follows:

1. We propose a multi-factorial data-driven approach to the prediction of individual daily stress;
2. We validate our approach with a seven-months dataset collected from 111 subjects;
3. We provide a comprehensive analysis of the predictive power of the proposed approach and a comparison with approaches based only on single families of features (personality, weather conditions, mobile phone features) or pairwise combinations thereof (personality and weather conditions, personality and mobile phone features, weather conditions and mobile phone features).

2. RELATED WORK

A large body of research on stress detection focused on physiological measurements to infer stress levels (see [23], [34], [39]). Heart-rate variability, galvanic skin response, respiration, muscle activity and temperature are among the most relevant features. However, despite providing reliable insights on stress levels, this approach has major limitations because it comprises wearable sensors that need to be carried at all times to allow for continuous monitoring.

Among the different changes in physiological parameters that happen during stressful situations, variation in speech production has inspired a number of studies using acoustic sensing on smartphones. Research on stress detection based on voice analysis considered different speech characteristics such as pitch, glottal pulse, spectral slope and phonetic variations. For example, Lu and colleagues [32] proposed StressSense, an Android application for stress detection from human voice in real-life conversation, and they achieved 81% and 76% accuracy for indoor and outdoor environments.

However, these methods depend on sound quality, which is not granted in natural settings (e.g., crowded public places, noisy outdoor), and the correlation between speech and emotion is subjected to large individual differences [43]. Hence, our performance of 72.28% is a good and reliable alternative to stress detection. Other studies focused on the video analysis of behavioural correlates of psychological stress [18]. These systems, despite providing an unobtrusive method for stress monitoring, cannot be employed in a large variety of real world and mobile environments and pose privacy concerns related to the recording of people’s behaviour.

A promising approach that can overcome the major shortcomings of stress detection based on physiological measures and on audio/video analysis is activity recognition from smartphone usage patterns. Studies in this field have been mainly focused on the understanding of relational dynamics of individuals [14]. Recently studies have started to investigate how smartphone usage habits can provide insights into users’ affective state [31] and stress levels [1]. LiKamWa and colleagues [31] proposed MoodScope, a mobile software system

that recognizes the users' mood, but not stress states, from smartphone usage analysis. They collected usage data and self-reported mood in a two months longitudinal study and used them to train mood models. Smartphone usage data consisted in phone calls, SMSes, e-mail messages, application use, web browsing histories and location changes, while self-reported mood was collected from users' input at least four times a day. MoodScope reached a 66% accuracy of participants' daily-average mood, with phone calls and categorized applications as the most useful features for mood discrimination.

Bauer and Lukowicz [1] focused on mid-term stress detection, monitoring 7 students during a two week exam session followed by two weeks of non-stressful period. The recorded data were related to participants' mobility patterns and social interactions, and included users' location, Bluetooth proximity, phone calls and SMSes. These features allowed to detect an average behaviour modification of 53% for each user during the exam session. A limitation of this study is the small number of subjects. Our multifactorial approach outperforms the approach proposed by [1] although a direct comparison may be not adequate given the different focus: our approach tend to daily classify people as "not stressed" or "stressed", while Bauer and Lukowicz try to detect stressful situations.

In 2013, Sano and Picard [42] reported an accuracy performance in stress recognition of 75% using a combination of features obtained from mobile phones and wearable sensors. However, the limited number of subjects used in their experiments (18) and the limited number of days (5) make preliminary the results of this study.

3. DATA COLLECTION

From November 12, 2010 to May 21, 2011, we collected a dataset capturing the lives of 117 subjects living in a married graduate student residency of a major US university. Our sample of subjects has a large variety in terms of provenance and cultural background: we have subjects from 16 countries such as USA, China, Israel, India, Iran, Russia, etc. During this period, each participant was equipped with an Android-based cellular phone incorporating a sensing software explicitly designed for collecting mobile data. Such software runs in a passive manner and does not interfere with the every day usage of the phone. The data collected consisted of: (a) call logs, (b) sms logs, (c) proximity data, obtained by scanning near-by phones and other Bluetooth devices every five minutes, and (d) data from surveys administered to participants, which provided self-reported information about personality traits ("Big Five") and self reported information about daily stress.

Proximity interaction data were derived from Bluetooth hits in a similar way as in previous reality mining studies [13]. Bluetooth scans were performed every 5 minutes in order to keep the battery from draining while achieving a high enough temporal resolution. The Bluetooth log of a given smartphone were then used to extract the list of the other participants' phones which were in proximity.

In total, the dataset consisted of 33497 phone calls, 22587 SMS, and 1460939 Bluetooth hits.

3.1 Stress data

At the evening, the participants were also asked to fill daily surveys about their daily self-perceived stress level.

The stress information was reported by the participants filling a seven items scale with 1 = "not stressed", 4 = "neutral" and 7 = "extremely stressed". In our experiments we used the data only for the subjects (111 subjects) who had provided at least 2 weeks of consecutive data.

The distribution of daily stress is visualized in Fig. 1. We see that it has a small negative skew – the density is moved to the higher region of stress score. The distribution has negative excess kurtosis, which in our case means that the sample reported a specific daily stress score more often than the neutral. Fig. 2 shows that within-person daily stress variance is more spread than between-person, but the density of between-person variance is higher.




Figure 1: Recorded Stress Scores Density




Figure 2: Within- and between-subject variance

3.2 Personality Data

Several studies in social psychology investigated the relationships between personality traits and psychological stress. Personalities that tend to be more negative are usually associated with greater distress, while outgoing and positive

Table 1: Selected Features Ranked by Mean Decrease in Accuracy

Rank	Feature	0	1	Mean Decrease in Accuracy	Mean Decrease in Gini Index
1	personality.Conscientiousness	13.65	18.04	23.35	159.96
2	personality.Agreeableness	14.22	19.73	22.92	167.30
3	personality.Neuroticism	15.96	21.04	22.56	183.87
4	personality.Openness	14.20	14.18	21.38	139.23
5	personality.Extraversion	15.75	15.02	21.07	158.51
6	weather.MeanTemperature	14.50	6.34	17.44	322.27
7	sms.RepliedEvents.Latency.Median	8.83	13.85	15.63	48.74
8	weather.Humidity	15.33	2.10	15.45	298.13
9	sms.AllEventsPerDay	8.61	0.56	10.50	42.91
10	bluetooth.Q95TimeForWhichIdSeen	4.99	6.05	9.94	32.47
11	bluetooth.MaxTimeForWhichIdSeen	6.24	7.23	9.47	32.12
12	sms.IncomingAndOutgoingPerDay	7.45	1.26	9.38	41.59
13	weather.Visibility	9.94	1.26	9.22	251.27
14	weather.WindSpeed	8.77	1.30	8.67	282.10
15	bluetooth.Q90TimeForWhichIdSeen	4.24	6.75	8.64	28.41
16	bluetooth.TotalEntropyShannon	5.04	3.51	8.56	31.37
17	call.EntropyMillerMadowOutgoingTotal	4.25	4.10	8.54	27.49
18	call.EntropyShannonOutgoingAndIncomingTotal	4.23	4.86	8.53	26.28
19	bluetooth.TotalEntropyMillerMadow	5.06	4.22	8.50	32.09
20	bluetooth.IdsMoreThan09TimeSlotsSeen	6.11	5.85	8.43	27.88
21	bluetooth.IdsMoreThan04TimeSlotsSeen	6.34	4.59	8.04	24.64
22	call.EntropyShannonMissedOutgoingTotal	3.13	4.92	7.85	24.34
23	bluetooth.IdsMoreThan19TimeSlotsSeen	2.97	5.16	7.78	20.87
24	call.EntropyShannonOutgoingTotal	3.10	6.45	7.78	24.79
25	bluetooth.Q75TimeForWhichIdSeen	5.16	4.70	7.76	22.07
26	call.EntropyMillerMadowMissedOutgoingTotal	4.09	5.45	7.55	24.64
27	call.EntropyMillerMadowOutgoingAndIncomingTotal	3.87	6.29	7.51	28.63
28	sms.OutgoingAndIncomingTotalEntropyMillerMadow	4.68	3.84	7.19	17.63
29	sms.OutgoingTotalEntropyMillerMadow	5.22	1.49	7.19	18.88
30	bluetooth.Q50TimeForWhichIdSeen	1.53	7.29	7.08	18.91
31	bluetooth.Q68TimeForWhichIdSeen	2.36	5.96	6.68	19.05
32	sms.OutgoingTotalEntropyShannon	2.53	2.77	5.13	17.59

personalities generally experience less distress [50], [52]. The majority of the studies that have examined the relationship between personality and distress focused on the Big Five traits [25], a personality model owing its name to the five traits it takes as a constitutive of people’s personality: Extraversion, Neuroticism, Agreeableness, Conscientiousness, and Openness to experience. Researchers showed significant associations between psychological stress, on the one hand, and Neuroticism, Extraversion and Conscientiousness, on the other. Duggan et al. [12] found that individuals high in Neuroticism may be more vulnerable to experiencing distress as they respond more negatively to daily stressors and report more daily stressful events and higher levels of daily stress. When people with high scores in Neuroticism encounter stressful events, they tend to experience them as more aversive than those low in this trait [3], [21]. Finally, in a study with university students, Volrath and Torgensen [52] showed that students with more adaptive personalities such as Extraversion and Conscientiousness are more positive and sociable and hence less affected by daily stress.

In our study, Big Five personality traits were measured by asking subjects to answer the online version of the 44 questions Big Five questionnaire developed by John et al. [25], by means of 5-point likert scales. The scores on the five traits were the average over the raw scores (inverted when needed) of the items pertaining to each trait.

3.3 Weather data

The question about the relationship between mood, health and weather has been extensively debated [22], [41]. Studies in environmental psychology investigated the role of weather as a stressor and showed significant effects of temperature, hours of sunshine and humidity on mood [24], [41]. A large-

scale study by Faust and colleagues [15] on 16,000 students in Switzerland showed an association between weather and sleep disorders, depressed mood and irritability. More recently, Denissen et al. [10] investigated the effects of six daily weather variables (temperature, wind power, sunlight, precipitation, air pressure, photoperiod) on three mood variables (positive affect, negative affect, and tiredness). Their results revealed main effects of temperature, wind power, and sunlight on negative affect, while sunlight had also a main effect on tiredness and mediated the effects of precipitation and air pressure on tiredness.

In our experiments, we used the following weather variables: (i) mean temperature, (ii) pressure, (iii) total precipitation, (iv) humidity, (v) visibility and (vi) wind speed (measured in m/s) metrics. The source weather data were collected from Wolfram Alpha ². All weather metrics are computed on a daily scale for the same day that is under investigation and the source data are extracted from the Boston area weather stations (e.g. KBOS) located on the same relative elevation as the campus where the data collection was performed.

4. FEATURE EXTRACTION

Based on previous works that characterize social interactions by means of mobile phone data and use social interactions data to predict people’s behaviors, states [2, 31], and traits [6, 9, 37], we derived the 25 call and sms basic features reported in Table 2 and the 9 proximity basic features reported in Table 3.

For each basic feature, we calculated second order features, such as mean, median, min, max, 99%, 95% quan-

²<http://www.wolframalpha.com>

tiles, quantiles corresponding to 0.5, 1, 1.5 and 2 standard deviations (applying Chebyshev's inequality), variance and standard deviation functions. Moreover, for each basic feature we calculated the same functions as above for 2 and 3 days backward-moving window to account for the possibility that past events influenced the current stress state.

In the following subsections we will describe more in detail the 25 call and sms basic features and the 9 proximity basic features.

4.1 Call and Sms Features

The features reported in Table 2 fall under four broad categories: (i) general phone usage, (ii) active behaviors, (iii) regularity, and (iv) diversity.

Table 2: List of Basic Features

General Phone Usage	
1. Total Number of Calls (Outgoing+Incoming)	
2. Total Number of Incoming Calls	
3. Total Number of Outgoing Calls	
4. Total Number of Missed Calls	
5. Number of SMS received	
6. Number of SMS sent	
Diversity	
7. Number of Unique Contacts Called	
8. Number of Unique Contacts who Called	
9. Number of Unique Contacts Communicated with (Incoming+Outgoing)	
10. Number of Unique Contacts Associated with Missed Calls	
11. Entropy of Call Contacts	
12. Call Contacts to Interactions Ratio	
13. Number of Unique Contacts SMS received from	
14. Number of Unique Contacts SMS sent to	
15. Entropy of SMS Contacts	
16. Sms Contacts to Interactions Ratio	
Active Behaviors	
17. Percent Call During the Night	
18. Percent Call Initiated	
19. Sms response rate	
20. Sms response latency	
21. Percent SMS Initiated	
Regularity	
22. Average Inter-event Time for Calls (time elapsed between two events)	
23. Average Inter-event Time for SMS (time elapsed between two events)	
24. Variance Inter-event Time for Calls (time elapsed between two events)	
25. Variance Inter-event Time for SMS (time elapsed between two events)	

Features for *general phone usage* consist of: the total number of outgoing, incoming and missed calls and the total number of sent/received sms. Moreover, we also computed the following ratios: outgoing to incoming calls, missed to (outgoing + incoming) calls, sent to received sms.

Then, we captured the *active behaviors* of an individual computing the following features: (i) percentage of calls done during the night, (ii) percentage of initiated calls during the night, (iii) the sms response rate, (iv) the sms response latency, and (v) the percentage of initiated sms. In particular, we consider a text from a user (A) to be a response to a text received from another user (B) if it is sent within an hour after user A received the last text from user B. The response rate is the percentage of texts people respond to. The latency is the median time it takes people to answer a text. Note that by definition, latency will be less or equal to one hour.

Diversity and *regularity* have been shown to be important for the characterization of different facets of human behavior. In particular, entropy, used as a measure of diversity, has been successfully applied to predict mobility [47], spending patterns [28, 46], online behavior [44] and person-

ality traits [37]. Concerning *regularity* features, we measured the time elapsed between calls, the time elapsed between sms exchanges and the time elapsed between call and sms. More precisely, we consider both the average and the variance of the inter-event time of one's call, sms and sum thereof (call+sms). Noticeably, in fact, even when two users have the same inter-event time for both call and sms, that quantity can be different for their sum.

Diversity measures how evenly an individual's time is distributed among others. In our case, the diversity of user behavior is addressed by means of three kinds of features: (i) entropy of contacts, (ii) unique contacts to interactions ratio, (iii) number of unique contacts, all computed both on calls and on sms. In particular, the entropy of an individual is the ratio between his/her total number of contacts and the relative frequency at which he/she interacts with them. The more one interacts equally often with a large number of contacts, the higher the entropy will be. For entropy calculation, we applied *Miller-Madow correction* [35], which is explained in Equation 1.

$$\hat{H}_{MM}(\theta) \equiv - \sum_{i=1}^p \theta_{ML,i} \log \theta_{ML,i} + \frac{\hat{m} - 1}{2N}, \quad (1)$$

where \hat{m} is a number of bins with nonzero θ -probability. The likelihood function is given as the product of probability density functions $P(\theta) = f(x_1; \theta) f(x_2; \theta) \cdots f(x_n; \theta)$ for a random sample X_1, \dots, X_n . θ_{ML} is the maximum likelihood estimate of θ , which maximizes $P(\theta)$. Miller-Madow correction was applied, dealing with the data quality problems, to get bias-corrected empirical entropy estimate.

4.2 Proximity Features

Starting from the Bluetooth hits collected, we filtered out all the cases with $RSSI < 0$. From the filtered Bluetooth proximity data we extracted the following basic Bluetooth proximity features (Table 3). In this case, the extracted fea-

Table 3: List of Basic Bluetooth Proximity Features

General Bluetooth Proximity	
1. Number of Bluetooth IDs	
2. Times most common Bluetooth ID is seen	
3. Bluetooth IDs accounting for n% of IDs seen	
4. Bluetooth IDs seen for more than k time slots	
5. Time interval for which a Bluetooth ID is seen	
6. Entropy of Bluetooth contacts	
Diversity	
7. Contacts to interactions ratio	
Regularity	
8. Average Bluetooth interactions inter-event time (time elapsed between two events)	
9. Variance of the Bluetooth interactions inter-event time (time elapsed between two events)	

tures fall under three broad categories: (i) general proximity information, (ii) diversity, and (iii) regularity. As for call and sms, we applied Miller-Madow correction for entropy calculation.

5. METHODOLOGY

We formulated the automatic recognition of daily stress as a binary classification problem ("not stressed" vs "stressed"), with labels 0 for "not stressed" and label 1 for "stressed". The two classes included all the cases with scores ≤ 4 and scores > 4 , respectively. The sizes of the resulting two classes are

36.16% for "stressed" and 63.84% for "not stressed". The inclusion of the cases with stress=4 in the 0 class meant to provide a more clearcut distinction between the "stressed" and the "non-stressed" cases.

The data set was then randomly split into a training (80% of data) and a testing (20% of data) dataset, carefully avoiding that data for the same subjects appeared in both the training- and in the test-set. In order to accelerate the convergence of the models, we *normalized* each dimension of the feature vector [4]. Additionally, we also used a leave-one-subject-out cross-validation strategy. Hence, 111 models for each personality trait were trained on 110-subject subsets, evaluating them against the remaining ones and finally averaging the results. The results obtained are not significantly different from the ones obtained using the random split 80% vs 20%. In the rest of the paper, we will discuss only the results obtained with the random split 80% vs 20%.

5.1 Feature Selection

As an initial step, we carried out a *Pearson correlation analysis* to visualize and better understand the relations between variables in the feature space. We found quite a large subset of features with strong mutual correlations and another subset of uncorrelated features. Hence, there was room for feature space reduction. We excluded using *principal component analysis* (PCA) because the transformation it is based on produces new variables that are difficult to interpret in terms of the original ones making the interpretation of the results more complex.

Therefore, we turned to a pipelined *variable selection* approach, based on *feature ranking* and *feature subset selection*, which was performed using only data from the training set. The metric used for feature ranking was the mean decrease in the *Gini coefficient of inequality*. This choice was motivated because it outperformed other metrics, such as mutual information, information gain and chi-square statistic with an average improvement of approximately 28.5%, 19% and 9.2% respectively [45]. The Gini coefficient ranges between 0, expressing perfect equality in predictive power and 1, expressing maximal inequality in predictive power. The feature with maximum mean decrease in Gini coefficient is expected to have the maximum influence in minimizing the out-of-the-bag error. It is known in the literature that minimizing the out-of-the-bag error results in maximizing common performance metrics used to evaluate models (*e.g.* accuracy, F1, AUC, etc.) [51].

The feature selection procedure produced a reduced subset of 32 features from an initial pool of about 500 features. Hence, we obtained a low-dimensional feature space that makes our approach efficient to implement into mobile and multimedia applications.

5.2 Model Building

We trained a variety of classifiers: (i) an ensemble of tree classifiers based on a Random Forest algorithm [5], (ii) a Generalized Boosted Model (GBM) [16], (iii) Support Vector Machines with linear and Gaussian radial basis kernels, and (iv) Neural Networks. The best solution of the classification problem was found using an ensemble of tree classifiers based on *Random Forest* algorithm. In the rest of the paper, we report the performance results only for Random Forest.

Random forest algorithm produces a combination of simple decision tree predictors, such that each tree is dependent

on the values of a random vector sampled independently with the same distribution for all the classification trees in the forest [5]. The decision boundary is formed according to the margin function. Given an ensemble of tree classifiers $h_1(\mathbf{x}), h_2(\mathbf{x}), \dots, h_K(\mathbf{x})$ and if the training set is drawn at random from the empirical distribution of the random vector Y, \mathbf{X} the margin function is defined as:

$$mg(\mathbf{X}, Y) = avg_k I(h_k(\mathbf{X}) = Y) - \max_{j \neq Y} avg_k I(h_k(\mathbf{X}) = j), \quad (2)$$

where $I(\cdot)$ is the characteristic function. The margin function measures the distance between the average votes at (\mathbf{X}, Y) for the right class and the average vote for any other class. For this model the generalization error function is:

$$PE^* = P_{\mathbf{X}, Y}(mg(\mathbf{X}, Y) < 0), \quad (3)$$

where $P_{\mathbf{X}, Y}$ is the probability over $\langle \mathbf{X}, Y \rangle$ space. For any event $A \subset \Omega$ of the feature space the characteristic function $I(\cdot)$ of A is:

$$I_A(x) = \begin{cases} 1 & \iff (x \in A) \\ 0 & \text{otherwise} \end{cases} \quad \begin{cases} 1 & \iff \exists x \\ 0 & \text{otherwise} \end{cases} \quad (4)$$

Random Forests classifiers were trained with a stepwise increase of the number of trees equal to the upper limit of 2^{11} . Optimal number of trees for model generalization as measured by mean misclassification rate for 10-fold cross-validation strategy is estimated to be 112 trees.

In order to find the final model, we trained a number of models and selected the best one based on κ metrics for the 10-fold validation strategy. The Cohen's κ measures pairwise agreement among a set of functions which are making classification decisions with correction for an expected chance agreement [7]:

$$\kappa = \frac{P(A) - P(E)}{1 - P(E)} \quad (5)$$

$\kappa = 0$ if there is no agreement more than expected by chance following the empirical distribution; while $\kappa = 1$ when there is a *max* agreement. κ is a state-of-the-art statistics about how significantly the classification model is different from chance. Importantly, κ is a more robust measure than the simple percent agreement, given that it takes into account chance agreement occurring without being too conservative.

During the learning and model selection process we used a random sampling with replacement to generate a new set of data for each fold from the basic training set, and followed leave-one-out 10-fold cross validation scheme. We adopted this strategy in order to prevent data overfitting and to deal with potential data loss in cases where calls, sms and Bluetooth proximities existed in the real world but were not registered by the smartphone logger software. Our *structural risk minimization*, as opposed to empirical risk minimization solution to prevent data overfitting, was incorporated by working with a regularization penalty into the learning process, balancing the model's complexity against training data fitting and by sampling the model training sets in such a way that they mimic the empirical distributions without most probable erroneous outliers.

Model parameter estimation selection was done iteratively on the basis of our exploratory analysis, inferred knowledge of the relationships between variables and model performance metrics (κ and Accuracy). Confounding variables are

identified but not removed from the dataset during training and test phases.

6. EXPERIMENTAL RESULTS

The performance metrics used to evaluate our approach are: accuracy, κ , sensitivity, and specificity. The recognition model based on random forest algorithm shows 90.68% accuracy on the training set and 72.28% accuracy on the test set. In Table 4 we provide the final stress recognition model performances on the test set along with their statistical significance [17].

Metric	Value
Accuracy	0.7228
95% CI	(0.7051, 0.7399)
No Information Rate	0.6384
P-Value [Acc > NIR]	< 2.2e-16
Kappa	0.3752
Sensitivity	0.5272
Specificity	0.8335
'Positive' Class	"stressed"

Table 4: Recognition Model Performance Metrics

Information about accuracy and κ metrics distribution using 10-fold cross validation strategy is provided in Table 5. As we can see, the distribution of the estimated performance metrics does not vary substantially among folds, signaling a good generalization despite the possible existence of heterogeneous data in each fold and the "noise" coming from the resampling procedure.

	Accuracy	Kappa
Min.	0.6959	0.2995
1st Qu.	0.7156	0.3535
Median	0.7282	0.3817
Mean	0.7232	0.3684
3rd Qu.	0.7312	0.3869
Max.	0.7404	0.4010

Table 5: 10-fold Cross-Validation Metrics

We also compared our approach based on combining multiple indicators with simpler approaches using as predictors (i) only personality traits, (ii) only weather conditions, (iii) only activities inferred from mobile phone data, (iv) a combination of personality traits and weather conditions, (v) a combination of personality traits and activities inferred from mobile phone data, (vi) a combination of weather conditions and activities inferred from mobile phone data. Table 6 reports accuracy, κ , sensitivity, specificity and F1 for each approach. In this table we also report the performance of (vii) a simple majority classifier, which always returns the majority class as prediction (accuracy = 63.84%). Finally, we also ran experiments with three classes ("not stressed", "neutral", "stressed"), with labels -1 for "not stressed", label 0 for "neutral", and label 1 for "stressed". The class "not stressed" included all the cases with scores < 4, the class "neutral" included all the cases with scores = 4, and the class "stressed" included all the scores > 4. The sizes of the resulting three classes are 42.83% for "not stressed", 20.98%

for "neutral", and 36.15% for "stressed". The global accuracy obtained by our multifactorial model, 59.57%, significantly outperformed the performance of simple majority classifier, which always returns the class "not stressed" as prediction.

7. DISCUSSION

The comparison among the performance of the various models in Table 6 provides convincing evidence that none of the features sets (personality, weather, smartphone activity) considered alone is endowed with a good enough predictive power. This conclusion applies also to pairwise combinations of the same features sets to the extent that neither personality+smartphone activity, nor personality+weather, nor weather+smartphone activity do any better than the majority classifier (accuracy=63.84%). Interestingly, significant improvements over the latter can only be obtained by the simultaneous usage of the three features sets: our final model based on a Random Forest classifier using 32-dimensional feature vectors obtained a 72.28% accuracy for our 2-class classification problem.

As pointed out in Section 2, some recent works have used mobile phones data for stress recognition [1, 42]. Bauer and Lukowicz [1] reported a 53% of accuracy in detecting the transition from stressful periods (a two week exam session) to non-stressful periods (two weeks after the exam session). Our multifactorial approach outperforms the approach proposed by [1] although a direct comparison may be not adequate given the different task. More recently, Sano and Picard reported an accuracy performance of 75% using a combination of features from mobile phones and more obtrusive wearable sensors. However, the limited number of subjects (only 18) and the limited number of days (only 5) make the results preliminary. Other approaches used video and audio features for stress recognition [18, 32]. For instance, StressSense, an application for stress detection from human voice, achieved a 76% of accuracy in outdoor environments. However, this method depends on sound quality and it may pose privacy concerns for people perceiving voice recording and analysis as a threat to their privacy. Hence, our performance of 72.28 shows that the proposed multifactorial approach is a reliable and less obtrusive alternative.

An investigation of the most important predictors of daily stress reveals interesting associations. Table 1 reports the 32 features selected and used in our model ranked by their mean reduction in accuracy. All the personality traits contribute significantly in predicting the daily stress variable. These results are interesting because the previous studies in social psychology focused their analyses mainly on the associations between stress and Neuroticism, Extraversion and Conscientiousness. Instead, our work shows the important contribution played also by Agreeableness and Openness to Experience to the automatic classification of daily stress. Moreover, these results open us the possibility of creating a multi-step stochastic model in which we first estimate the personality and then we use those estimates as independent variables for the daily stress recognition problem. Our current approach uses self-reported information on personality and this strategy could be a limitation for scaling to larger sample of users. However, recent studies showed that personality traits may be reliable recognized from mobile phone data [6, 9, 37, 48].

With regard to weather, we found confirmation for the association between temperature and stress. Moreover, sig-

Table 6: Model Metrics Comparison for Feature Subsets

Model	Accuracy	Kappa	Sensitivity	Specificity	F1
Our Multifactorial Model	72.28	37.52	52.72	83.35	57.89
Baseline Majority Classifier	63.84	0.00	100.00	0.00	0.00
Weather Only	36.16	0.00	100.00	0.00	0.00
Personality Only	36.16	0.00	100.00	0.00	0.00
Bluetooth+Call+Sms	48.59	6.80	73.80	34.32	50.94
Personality+Weather	43.55	2.96	81.90	21.83	51.20
Personality+Bluetooth+Call+Sms	46.40	7.01	83.17	25.57	52.88
Weather+Bluetooth+Call+Sms	49.60	-5.45	38.45	55.91	35.55

nificant effects of other meteorological variables – humidity, visibility, and wind speed – for predicting daily stress were also found.

Regarding the mobile phone data, it is interesting to note the contribution of proximity features. Out of the selected 32 features, 11 features are proximity ones, 6 comes from call data and 6 from sms data. In particular, an interesting predictive role is played by the number of time intervals for which an id is seen. The results obtained using proximity features seem to confirm previous findings in social psychology: in particular, the relevant role played by face-to-face interactions and by interactions with strong ties in determining the stress level of a subject [27]. For sure, this result requires further investigation. In addition, two features capturing the entropy in proximity interactions are among the selected ones. This finding seems to confirm results available in the social psychological literature about the associations between stress and the richness/diversity of social interactions [20]. Further confirmation to this conclusion comes from the similarly important role played by entropy-based call and sms features.

The remaining selected features related to sms interactions are (i) the latency in replying to a text message, defined as the median time to answer a text message and (ii) the amount of sms communications (outgoing+incoming).

8. IMPLICATIONS AND LIMITATIONS

Stress has become a major problem in our society. Ubiquitous connectivity, information overload, increased mental workload and time pressure are all elements contributing to increase general stress levels. While in some cases people may realize that they are undergoing stressful situations, severe and chronic stress may be more difficult to detect. Moreover, stress may be considered the norm in a modern and demanding society. Nonetheless, while slightly increased stress levels may be functional for productivity, prolonged and severe stress can be at the source of several physical dysfunctions like headache, sleep or immunological disorders, unhealthy behaviours such as smoking and bad eating habits, as well as of psychological and relational problems. Beside manifest social costs, stress also entails considerable financial costs for our economies, which are estimated by the World Health Organization in 300 billion dollars a year for American enterprises, and 20 billion euro for Europe ones, in terms of absenteeism and low productivity.

Our technology provides a cost-effective, unobtrusive, widely available and reliable tool for stress recognition. It detects daily stress levels with a 72.28% accuracy combining real life data from different sources, such as personality traits, social relationships (in terms of calls, sms and proximity interactions), and weather data. The development of a reliable stress recognition system is a first but essential step

toward wellbeing and sustainable living, and its scope can be extended to different areas of applicability. Providing people with a tool capable of gathering rich data about real life, and transforming them into meaningful insights about stress levels, paves the way to a new generation of context-aware technologies that can target therapists, enterprises and common citizens.

This technology can inform the design of automatic systems for the assessment and treatment of psychological stress. With such a tool, therapists could monitor and record patients' daily stress levels, access longitudinal data, identify recurrent or significant stressors and modulate treatment accordingly.

In work environments, where stress has become a serious problem affecting productivity, leading to occupational issues and causing health diseases, our system could be extended and employed for early detection of stress-related conflicts and stress contagion, and for supporting balanced workloads. Awareness is a first but crucial step to motivate people to change their behaviour and take informed and concrete steps toward a healthy lifestyle and appropriate stress coping strategies. Mobile applications developed on the basis of our technology could provide feedback to increase people's awareness of their stress levels, alerts when they reach a warning threshold, and suggest stress management and relaxation techniques when appropriate.

However, our study has also some limitations. We can list the following ones: (i) our sample comes from a population living in the same environment. Our subjects were all married graduate students living in a campus facility of a major US university; and (ii) the non-availability of proximity data concerning the interaction with people not participating in the data collection, a fact that is common to many other relevant studies and that has been also pointed out by [40]. The first problem is at least partially attenuated by the large variability of the sample in terms of provenance and cultural background (in our sample we have subjects from 16 countries and from all the continents), which can be expected to correspond to a wide palette of interaction behaviors that efficaciously counterbalance the effects of living-place homogeneity.

9. CONCLUSION

The goal of this paper was to investigate the automatic recognition of people's daily stress from three different sets of data: a) people activity, as detected through their smartphones (data pertaining to transitory properties of individuals); b) weather conditions (data pertaining to transitory properties of the environment); and c) personality traits (data concerning permanent dispositions of individuals). The problem was modeled as a 2-way classification one. The results convincingly suggest that all the three

types of data are necessary for attaining a reasonable predictive power. As long as one of those information sources is dropped, performances drop below those of the baselines. Moreover, the distributional data for accuracy and κ show the robustness and generalization power of our multifactorial approach.

Taken together, and despite the limitations discussed above, our results not only provide evidence that individual daily stress can be reliably predicted, but they also point to the necessity of considering at the same time people's transitory properties (smartphone activity), transitory properties of the environment and information about stable individual characteristics. For the sake of transitory individual properties, mobile phone usage patterns have important advantage over alternative methods: they are less unobtrusive and raise limited privacy problems as compared to, e.g., voice analysis or the exploitation of data from physiological sensors. Moreover, and importantly, automatic stress detection based on mobile phone data can take advantage of the extensive usage and diffusion of these devices, it can be applied in several real world situations and it can be exploited for a variety of applications that are delivered by means of the same device. For example, applications used to inform the design of clinical decision support systems or self-monitoring applications of stress levels in work settings and in other daily life situations, which allows people to identify personal stressors and enforces their proactive role in stress prevention and management.

Acknowledgment

This work was partially supported by Telecom Italia Semantics and Knowledge Innovation Laboratory (SKIL) with research grant T.

10. REFERENCES

- [1] G. Bauer and P. Lukowicz. Can smartphones detect stress-related changes in the behaviour of individuals? In *Pervasive Computing and Communications Workshops (PERCOM Workshops)*, 2012 *IEEE International Conference on*, pages 423–426. IEEE, 2012.
- [2] A. Bogomolov, B. Lepri, and F. Pianesi. Happiness recognition from mobile phone data. *Proceedings of the 2013 International Conference on Social Computing (SocialCom 2013)*, pages 790–795, 2013.
- [3] N. Bolger and E. A. Schilling. Personality and the problems of everyday life: The role of neuroticism in exposure and reactivity to daily stressors. *Journal of Personality*, 59(3):355–386, 1991.
- [4] G. E. P. Box and D. R. Cox. An Analysis of Transformations. *Journal of the Royal Statistical Society. Series B (Methodological)*, 26(2):211–252, 1964.
- [5] L. Breiman. Random forests. *Mach. Learn.*, 45(1):5–32, Oct. 2001.
- [6] G. Chittaranjan, J. Blom, and D. Gatica-Perez. Mining large-scale smartphone data for personality studies. *Personal Ubiquitous Comput.*, 17(3):433–450, Mar. 2013.
- [7] J. Cohen. A Coefficient of Agreement for Nominal Scales. *Educational and Psychological Measurement*, 20(1):37–46, Apr. 1960.
- [8] S. Cohen, K. R. C., and L. U. Gordon. *Measuring stress: A guide for health and social scientists*. Oxford University Press, USA, 1997.
- [9] R. de Oliveira, A. Karatzoglou, P. Concejero Cerezo, A. Armenta Lopez de Vicuña, and N. Oliver. Towards a psychographic user model from mobile phone usage. In *CHI'11 Extended Abstracts on Human Factors in Computing Systems*, pages 2191–2196. ACM, 2011.
- [10] J. Denissen, L. Butalid, L. Penke, and M. Van Aken. The effects of weather on daily mood: A multilevel approach. *Emotion Researcher*, 8(5):662–667, 2008.
- [11] W. Dong, B. Lepri, and A. Pentland. Modeling the co-evolution of behaviors and social relationships using mobile phone data. In *MUM*, pages 134–143, 2011.
- [12] C. Duggan, P. Sham, A. Lee, C. Minne, and R. Murray. Neuroticism: a vulnerability marker for depression evidence from a family study. *Journal of Affective Disorders*, 35(3):139 – 143, 1995.
- [13] N. Eagle and A. Pentland. Reality mining: sensing complex social systems. *Personal and Ubiquitous Computing*, 10(4):255–268, 2006.
- [14] N. Eagle, A. S. Pentland, and D. Lazer. Inferring friendship network structure by using mobile phone data. *Proceedings of the National Academy of Sciences*, 106(36):15274–15278, 2009.
- [15] V. Faust, M. Weidmann, and W. Wehner. The influence of meteorological factors on children and youths: A 10% random selection of 16,000 pupils and apprentices of basle city (switzerland). *Acta Paedopsychiatrica: International Journal of Child & Adolescent Psychiatry*, 1974.
- [16] Y. Freund and R. E. Schapire. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences*, 55(1):119 – 139, 1997.
- [17] S. Garcia and F. Herrera. An Extension on "Statistical Comparisons of Classifiers over Multiple Data Sets" for all Pairwise Comparisons. *Journal of Machine Learning Research*, 9:2677–2694, Dec. 2008.
- [18] D. Giakoumis, A. Drosou, P. Cipresso, D. Tzovaras, G. Hassapis, A. Gaggioli, and G. Riva. Real-time monitoring of behavioural parameters related to psychological stress. *Studies in health technology and informatics*, 181:287, 2012.
- [19] M. Gonzalez, C. Hidalgo, and L. Barabasi. Understanding individual mobility patterns. *Nature*, 453(7196):779–782, 2008.
- [20] W. J. Goode. A theory of role strain. *American Sociological Review*, 25(4):pp. 483–496, 1960.
- [21] L. H. A. S. Gunthert, Kathleen Cimboic, Cohen. The role of neuroticism in daily stress and coping. *Journal of Personality and Social Psychology*, 77:1087–1100, 1999.
- [22] J. Hardt and H. Gerbershagen. No changes in mood with the seasons: observations in 3000 chronic pain patients. *Acta Psychiatrica Scandinavica*, 100(4):288–294, 1999.
- [23] J. Healey and R. Picard. Detecting stress during real-world driving tasks using physiological sensors. *Intelligent Transportation Systems, IEEE Transactions on*, 6(2):156–166, 2005.

- [24] E. Howarth and M. S. Hoffman. A multidimensional approach to the relationship between mood and weather. *British Journal of Psychology*, 75(1):15–23, 1984.
- [25] O. P. John and S. Srivastava. The Big Five trait taxonomy: History, measurement, and theoretical perspectives. In L. A. Pervin and O. P. John, editors, *Handbook of Personality: Theory and Research*, pages 102–138. Guilford Press, New York, second edition, 1999.
- [26] E. Jovanov, A. O'Donnell Lords, D. Raskovic, P. Cox, R. Adhami, and F. Andrasik. Stress monitoring using a distributed wireless intelligent sensor system. *Engineering in Medicine and Biology Magazine, IEEE*, 22(3):49–55, 2003.
- [27] D. Krackhardt. The strength of strong ties: The importance of philos in organizations. pages 216–239, 1992.
- [28] C. Krumme, A. Llorente, M. Cebrian, A. Pentland, and E. Moro. The predictability of consumer visitation patterns. *Scientific Reports*, (1645), 2013.
- [29] N. D. Lane, E. Miluzzo, H. Lu, D. Peebles, T. Choudhury, and A. T. Campbell. A survey of mobile phone sensing. *IEEE Communications Magazine*, 48(9):140–150, 2010.
- [30] N. Lathia, K. K. Rachuri, C. Mascolo, and P. J. Rentfrow. Contextual dissonance: Design bias in sensor-based experience sampling methods. *Proceedings of Ubicomp 2013*, pages 183–192, 2013.
- [31] R. LiKamWa, Y. Liu, N. D. Lane, and L. Zhong. Moodscope: building a mood sensor from smartphone usage patterns. In *Proceeding of the 11th annual international conference on Mobile systems, applications, and services, MobiSys '13*, pages 389–402, New York, NY, USA, 2013. ACM.
- [32] H. Lu, D. Fraundorfer, M. Rabbi, M. S. Mast, G. T. Chittaranjan, A. T. Campbell, D. Gatica-Perez, and T. Choudhury. Stresssense: detecting stress in unconstrained acoustic environments using smartphones. In *Proceedings of the 2012 ACM Conference on Ubiquitous Computing, UbiComp '12*, pages 351–360, New York, NY, USA, 2012. ACM.
- [33] A. Madan, M. Cebrián, S. T. Moturu, K. Farrahi, and A. Pentland. Sensing the "health state" of a community. *IEEE Pervasive Computing*, 11(4):36–45, 2012.
- [34] D. Majoe, P. Bonhof, T. Kaegi-Trachsel, J. Gutknecht, and L. Widmer. Stress and sleep quality estimation from a smart wearable sensor. In *Pervasive Computing and Applications (ICPCA), 2010 5th International Conference on*, pages 14–19, 2010.
- [35] G. Miller. Note on the bias of information estimates. *Information theory in psychology: Problems and methods*, 2(95):100, 1955.
- [36] G. Miller. The smartphone psychology manifesto. *Perspectives on Psychological Science*, 7(3):221–237, 2012.
- [37] Y.-A. Montjoye, J. Quoidbach, F. Robic, and A. Pentland. Predicting personality using novel mobile phone-based metrics. In A. Greenberg, W. Kennedy, and N. Bos, editors, *Social Computing, Behavioral-Cultural Modeling and Prediction*, volume 7812 of *Lecture Notes in Computer Science*, pages 48–55. Springer Berlin Heidelberg, 2013.
- [38] A. Muaremi, B. Arnrich, and G. Tröster. A survey on measuring happiness with smart phones. In *6th International Workshop on Ubiquitous Health and Wellness (UbiHealth 2012)*, 2012.
- [39] K. Plarre, A. Raij, S. Hossain, A. Ali, M. Nakajima, M. Al'absi, E. Ertin, T. Kamarck, S. Kumar, M. Scott, D. Siewiorek, A. Smailagic, and L. Wittmers. Continuous inference of psychological stress from sensory measurements collected in the natural environment. In *Information Processing in Sensor Networks (IPSN), 2011 10th International Conference on*, pages 97–108, 2011.
- [40] D. Quercia, J. Ellis, L. Capra, and J. Crowcroft. Tracking "gross community happiness" from tweets. In *CSCW*, pages 965–968, 2012.
- [41] J. L. Sanders and M. S. Brizzolara. Relationships between weather and mood. *The Journal of General Psychology*, 107(1):155–156, 1982.
- [42] A. Sano and R. W. Picard. Stress recognition using wearable sensors and mobile phones. *Proceedings of the 2013 Humaine Association Conference on Affective Computing and Intelligent Interaction*, pages 671–676, 2013.
- [43] K. R. Scherer, D. Grandjean, T. Johnstone, G. Klasmeier, and T. Bänziger. Acoustic correlates of task load and stress. In *INTERSPEECH*, 2002.
- [44] R. Sinatra and M. Szell. Entropy and the predictability of online life. *Entropy*, (16):543–556, 2014.
- [45] S. R. Singh, H. A. Murthy, and T. A. Gonsalves. Feature selection for text classification based on gini coefficient of inequality. *Journal of Machine Learning Research-Proceedings Track*, 10:76–85, 2010.
- [46] V. K. Singh, L. Freeman, B. Lepri, and A. S. Pentland. Predicting spending behavior using socio-mobile features. In *Social Computing (SocialCom), 2013 International Conference on*, pages 174–179. IEEE, 2013.
- [47] C. Song, Z. Qu, N. Blumm, and A. Barabasi. Limits of predictability in human mobility. *Science*, (327):1018–1021, 2010.
- [48] J. Staiano, B. Lepri, N. Aharony, F. Pianesi, N. Sebe, and A. Pentland. Friends don't lie: inferring personality traits from social network structure. In *Proceedings of the 2012 ACM Conference on Ubiquitous Computing*, pages 321–330. ACM, 2012.
- [49] A. Stopczynski, V. Sekara, P. Sapiezynski, A. Cuttone, M. M. Madsen, J. E. Larsen, and S. Lehmann. Measuring large-scale social networks with high resolution. *PLoS ONE*, (9), 2014.
- [50] J. Suls, P. Green, and S. Hillis. Emotional reactivity to everyday problems, affective inertia, and neuroticism. *Personality and Social Psychology Bulletin*, 24(2):127–136, 1998.
- [51] E. Tuv, A. Borisov, G. Runger, and K. Torkkola. Feature selection with ensembles, artificial variables, and redundancy elimination. *The Journal of Machine Learning Research*, 10:1341–1366, 2009.
- [52] M. Vollrath and S. Torgersen. Personality types and coping. *Personal. Individ. Differ.*, 29:367–378, 2000.